

Introducing ICBe: Very High Recall and Precision Event Extraction from Narratives about International Crises

Rex W. Douglass^{a,1,2}, Thomas Leo Scherer^a, J. Andrés Gannon^b, Erik Gartzke^a, Jon Lindsay^c, Shannon Carcelli^d, Jonathan Wilkenfeld^d, David M. Quinn^e, Catherine Aiken^f, Jose Miguel Cabezas Navarro^g, Neil Lund^d, Egle Murauskaitė^h, and Diana Partridge^h

^aDepartment of Political Science, University of California, San Diego, CA, USA.; ^bBelfer Center for Science and International Affairs, Harvard Kennedy School, MA, USA.; ^cSchool of Cybersecurity and Privacy | Sam Nunn School of International Affairs, Georgia Institute of Technology, GA, USA.; ^dDepartment of Government and Politics, University of Maryland, College Park, MD, USA.; ^eFaculty Specialist, National Consortium for the Study of Terrorism and Responses to Terrorism (START), University of Maryland, College Park, MD, USA.; ^fEdmund A. Walsh School of Foreign Service, Georgetown University, Washington, DC, USA.; ^gSociety and Health Research Center, Universidad Mayor, Santiago, Chile.; ^hICONS Project // START, University of Maryland, College Park, MD, USA.; ⁱICONS Project // START, University of Maryland, College Park, MD, USA.

This manuscript was compiled on July 16, 2022

How do international crises unfold? We conceptualize of international relations as a strategic chess game between adversaries and develop a systematic way to measure pieces, moves, and gambits accurately and consistently over a hundred years of history. We introduce a new ontology and dataset of international events called ICBe based on very high-quality corpus of narratives from the International Crisis Behavior (ICB) Project. We demonstrate that ICBe has higher coverage, recall, and precision than existing state of the art datasets and conduct two detailed case studies of the Cuban Missile Crisis (1962) and Crimea-Donbas Crisis (2014). We further introduce two new event visualizations (event iconography and crisis maps), an automated benchmark for measuring event recall using natural language processing (synthetic narratives), and an ontology reconstruction task for objectively measuring event precision. We make the data, replication material, and visualizations of every historical episode available at a companion website www.crisisevents.org.

Diplomacy | War | Crises | International Affairs | Computational Social Science

If we could record every important interaction between countries in all of diplomacy, military conflict, and international political economy, how much unique information would this chronicle amount to, and how surprised would we be to see something new? In other words, what is the entropy of international relations? This record could in principle be unbounded, but the central conceit of social science is that there are structural regularities that limit what actors can do, their best options, and even which actors are likely to survive (1, 2). If so, then these events can be systematically measured, and accordingly, massive effort is expended in social science attempting to record these regularities.* With improvements in natural language processing, more open-ended efforts have sought to capture entire unstructured streams of international events drawn from news reports.[†] How close these efforts are to accurately measuring all or even most of what is essential in international relations is an open empirical question, one for which we provide new evidence here.

*See work on crises (3, 4), militarized disputes (5–7), wars (8, 9), organized violence (10, 11), political violence (12), sanctions (13), trade (14), and international agreements (15–17), dispute resolution (17, 18), and diplomacy (19, 20).

[†]See (21); (22); (23); (24); (25); (26). On event-extraction from images and social-media see (27) and (28).

Our contribution is a high coverage ontology and event dataset for key historical episodes in 20th and 21st-century international relations. We develop a large, flexible ontology of international events with the help of both human coders and natural language processing. We apply it sentence-by-sentence to an unusually high-quality corpus of historical narratives of international crises (1, 29–32). The result is a new lower bound estimate of how much actually happens between states during pivotal historical episodes. We then develop several methods for objectively gauging how well these event codings reconstruct the information contained in the original narrative. We conclude by benchmarking our event codings against several current state-of-the-art event data collection efforts. We find that existing systems produce sequences of events that do not contain enough information to reconstruct the underlying historical episode. The underlying fine-grained variation in international affairs is unrecognizable through the lens of current quantification efforts.

This is a measurement paper that makes the following argument — there is a real-world unobserved latent concept

Significance Statement

Countries routinely face crises that risk escalating into full scale war but we do not have systematic measurements of the progression of past crises and what moves and counter moves led to or helped avoid war. Instead policy makers typically rely on one or two historical analogies, chosen through ad hoc selection criteria, and described in unsystematic terms. This paper introduces a new scientific approach to measuring the step by step moves of international crises over the last hundred years, combining subject expertise with state of the art natural language processing and machine learning methods. It serves as a guide for constructing and evaluating large scale measurement collection in the social sciences.

Conceptualization: R.W.D., E.G., J.L.; Methodology: R.W.D., T.L.S.; Software: R.W.D.; Validation: R.W.D., T.L.S.; Formal Analysis: R.W.D., T.L.S.; Investigation: S.C., R.W.D., J.A.G., C.K., N.L., E.M., J.M.C.N., D.P., D.Q., J.W.; Data Curation: R.W.D., D.Q., T.L.S., J.W.; Writing - Original Draft: R.W.D., T.L.S.; Writing - Review & Editing: R.W.D., J.A.G., E.G., T.L.S.; Visualization: R.W.D., T.L.S.; Supervision: E.G.; Project Administration: S.C., R.W.D., J.A.G., D.Q., T.L.S., J.W.; Funding Acquisition: E.G., J.L.

Please declare any conflict of interest here.

² To whom correspondence should be addressed. E-mail: rex.douglass@gmail.com

known as international relations, we propose a method for systematically measuring it, we successfully apply this method producing a new large scale set of measurements, those measurements exhibit several desirable kinds of internal and external validity, and those measurements out-perform other existing approaches. The article organizes that argument into eight sections: task definition; corpus; priors/existing state of the art; ICBe coding process; internal consistency; case study selection; recall; and precision. A final section concludes.

Task Definition

We consider the measurement task of abstracting discrete events about a historical episode in international relations. The easiest way to convey the task is with an example. Figure 1 shows a narrative account of the Cuban Missile Crisis (1962) alongside a mapping from each natural language sentence to discrete machine readable abstractive events. Formally, a historical episode, H , is demarcated by a period of time $[T_{start}, T_{end}] \in T$, a set of Players $p \in P$, and a set of behaviors they undertook during that time $b \in B$. International Relations, IR , is the system of regularities that govern the strategic interactions that world actors make during a historical episode, given their available options, preferences, beliefs, and expectations of choices made by others. We observe neither H nor IR directly. Rather the Historical Record, HR , produces documents $d \in D$ containing some relevant and true (as well as irrelevant and untrue) information about behaviors that were undertaken recorded in the form of unstructured natural language text. The task is to combine informative priors about IR with an unstructured corpus D to produce a series of structured discrete events, $e \in E$, that have high coverage, precision, and recall over what actually took place in history, H .

Fig. 1. Case Study 1: Cuban Missile Crisis (1962) - ICB Narrative vs. ICBe Events

Natural Language Sentences (ICB Corpus)	Machine Readable Events (ICBe)
1 There were three crisis actors in the most ominous of all East/West crises, the United States, Cuba, and the Soviet Union.	
2 The crisis lasted from 16 October to 20 November 1962.	
3 After the 1961 Bay of Pigs abortive invasion (see Case #181), Cuba became one of the central issues of U.S. foreign policy: the U.S. viewed Cuba as a potential source of communist-oriented subversive activities in Latin America.	
4 When the U.S. discovered the presence of Soviet military personnel in Cuba on 7 September 1962 it called up 150,000 reservists.	USA mobilization 🚶 100ks
5 The Soviets mobilized on the 11th.	USA discover fact; deployment to area 🚶 100s; 🚶 1ks
6 Although persistent rumors circulated concerning the deployment of Soviet missiles in Cuba, Soviet Ambassador Anatoly Dobrynin denied the charges, and Premier Khrushchev gave his personal assurances that ground-to-ground missiles would never be shipped to Cuba.	
7 On the eve of the Missile crisis Washington did not openly challenge the Soviet statements concerning the defensive character of the weapons being sent to Cuba.	
8 The U.S. crisis was triggered on 16 October when the CIA presented to President Kennedy photographic evidence of the presence of Soviet missiles in Cuba.	USA discover fact; start of crisis - USSR deployment to area fortify 🚶 100s; 🚶 1ks
9 The U.S. responded with a decision on the 20th to blockade all offensive military equipment en route to Cuba.	USA start of crisis - USSR blockade 🚶 10ks
10 When this was announced on 22 October, a crisis was triggered for Cuba and the USSR. An urgent meeting of the UN Security Council was requested by both the U.S. and Cuba on the 22nd, and by the USSR the next day.	USA blockade 🚶 10ks; USSR start of crisis - USA
11 On the 23rd as well, the Soviets accused the United States of violating the UN Charter and announced an alert of its armed forces and those of the Warsaw Pact members.	USSR accuse USA - USA violate terms of treaty 🚶 100ks
12 That day Cuba responded by condemning the U.S. blockade and declaring its willingness to fight.	USSR disapprove USA
13 A resolution was adopted on the 23rd by the OAS calling for the withdrawal of the missiles from Cuba and recommending that member-states take all measures, including the use of force, to ensure that the government of Cuba would not continue to receive military material.	OAS demand USSR withdraw from area 🚶 1ks
14 On 24 October the Security Council adopted a resolution requesting the Secretary-General to confer with the parties.	UN express intent USSR; USA; mediation 🚶 10ks
15 On that same day, U Thant began mediation by sending identical letters to Khrushchev and Kennedy which proposed that the Soviet Union and the United States enter into negotiations, during which period both the shipment of arms and the quarantine would be suspended.	UN express intent - UN mediation 🚶 10ks
16 Moscow's major response to the crisis was a letter from Khrushchev to Kennedy on 26 October offering the removal of Soviet offensive weapons from Cuba and the cessation of further shipments in exchange for an end to the U.S. quarantine and a U.S. assurance that it would not invade Cuba.	USSR offer USA happens - USSR withdraw from area 🚶 100s; 🚶 1ks
17 The situation was exacerbated on the 27th when a U.S. U-2 surveillance plane was shot down.	USSR end blockade 🚶 100s
18 That day another Khrushchev letter was received in Washington offering the removal of Soviet missiles from Cuba in exchange for the removal of U.S. missiles from Turkey.	USSR offer USA happens - USA withdraw from area 🚶 100s; 🚶 1ks
19 U.S. mobilization and aerial reconnaissance flights were stepped up.	USA offer USA happens - USA withdraw from area 🚶 100s; 🚶 1ks
20 And on the 27th President Kennedy sent the Soviet premier an acceptance of the proposals contained in the letter of 26 October while making no reference to Khrushchev's second letter of the 27th.	UN accept USSR - USA
21 The following day Khrushchev notified the U.S. government that he had ordered work on the missile sites in Cuba stopped.	USSR accept USA - USA
22 He agreed to ship the missiles back to the USSR and promised that UN observers would be allowed to verify the dismantling of the sites.	
23 At the same time he warned Washington that U-2 reconnaissance flights over Cuba must be stopped as well.	USSR demand USA - USA withdraw from area 🚶 100s
24 The crisis continued at a lower level of intensity for several more weeks due to Cuban President Castro's demands concerning a U.S. pledge not to invade his country.	USSR de-mobilization; lower alert USA; body of water; coastline 🚶 100ks
25 On 30 October U Thant began talks in Havana, and Kennedy agreed to lift the quarantine for the duration of the talks.	UN end blockade 🚶 10ks
26 When Cuba rejected UN inspection, the U.S. resumed the quarantine and air surveillance.	UN accept USSR - USA end blockade 🚶 10ks
27 The Kremlin sent Deputy Premier Anastas Mikoyan to Cuba on 2 November to try to persuade Castro to allow UN inspection.	USSR reject USA; USA unspecifed cooperation 🚶 100s
28 When this proved unsuccessful, a U.S.-USSR agreement was reached on 7 November allowing U.S. inspection and interception of Soviet ships leaving Cuba and the photographing of the missiles.	UN discussion; meeting USSR
29 The following day the superpowers negotiated the removal of the IL-28 bombers which Castro had claimed were Cuban property.	USA sign formal agreement USA
30 Castro's agreement was conveyed to the U.S. on 20 November 1962, which terminated the Missile crisis for all three actors.	USSR withdraw from area 🚶 100s
31 The U.S. naval quarantine was lifted immediately, but aerial surveillance continued until the agreement was completely carried out.	USA sign formal agreement USA
32 U Thant's mediation efforts had an important impact on crisis resolution.	USSR end of crisis 🚶 100s
33 The acting Secretary-General managed to act as a facilitator and kept lines of communication open by contacting both parties.	UN end blockade 🚶 100s
34 He also used formulative mediation by highlighting common interests and making substantive suggestions to both sides for diffusing the crisis.	
35 Overall, U Thant's mediation effort was a significant factor in easing tensions between crisis actors which contributed to a more rapid termination of the crisis.	
36 Three regional/security organizations were involved in this crisis--the OAS, NATO, and the WTO. Kennedy sent Dean Acheson to Paris on 23 October to brief NATO's Permanent Council on U.S. measures against Cuba.	
37 Involved were members of the Warsaw Pact, whose forces were put on alert, and Latin American states which offered military assistance to the U.S.	

Corpus

For our corpus, D , we select a set of unusually high-quality historical narratives from the International Crisis Behavior (ICB) project ($n = 471$) (SI Appendix, Table A1) (33, 34). Their domain is 20th and 21st-century crises, defined as a change in the type, or an increase in the intensity, of disruptive interaction with a heightened probability of military hostilities that destabilizes states' relationships or challenges the structure of the international system (3).[†] Crises are a significant focus of detailed single case studies and case comparisons because they provide an opportunity to examine behaviors in IR short of, or at least prior to, full conflict (3, 35–42). Case selection was exhaustive based on a survey of world news archives and region experts, cross-checked against other databases of war and conflict, and non-English sources (33, 43). Each narrative was written by consensus by a small number of scholars, using a uniform coding scheme, with similar specificity (44). The corpus is unique in IR because it is designed to be used in a downstream quantitative coding project.

Prior Beliefs about IR, Ontological Coverage, and the Existing State of the Art

Next we draw informative prior beliefs about the underlying process of IR that we expect to govern behavior during historical episodes and their conversion to the historical record. We organize our prior beliefs along two overarching axes, summarized in detail by Table 1.

The first axis (rows) represents the types of information we expect to find in IR and forms the basis for our proposed ontology. We employ a metaphor of a chess game, with players (polities, rebel groups, IGOs, etc.), pieces (military platforms, civilians, domains), and behaviors (think, say, do). Precise sequencing is required to capture gambits (sequences of moves) and outcomes (victory, defeat, peace, etc.), while precise geocoding is required to understand the chessboard (medium of conflict). We find 464 actors and 117 different behaviors, and provide a full codebook in the online material.[§]

We base our informed priors primarily on two sources of information. The first is the extensive existing measurement efforts of IR which we provide citations to alongside each concept. Second, we performed preliminary natural language processing of the corpus and identified named entities and behaviors mentioned in the text. Verbs were matched to the most likely definition found in Wordnet (45), tallied, and then aggregated into a smaller number hypernyms balancing conceptual detail and manageable sparsity for human coding (SI Appendix, Table A2).

The second axis (columns) compares the very high ontological coverage of ICBe to existing state of the art systems in production and with global coverage. They begin with our contribution ICBe, alongside other event-level datasets including CAMEO dictionary lookup based systems (Historical Phoenix (46); ICEWS (24, 25); Terrier (26)), the Militarized Interstate Disputes Incidents dataset, and the UCDP-GED dataset (10, 11, 47).[¶] The final set of columns compares episode-level

datasets beginning with the original ICB project (3, 4, 56); the Militarized Interstate Disputes dataset (5, 6, 57, 58), and the Correlates of War (8). With the exception of large scale CAMEO dictionary based systems, the existing state of the art quantitative datasets ignore the vast majority of the information content found in international relations.^{||}

[†]On near crises see (32).

[§]See the Github Repository [ICBEEventData](#).

[¶]Additional relevant but dated or too small of an overlap in domain include BCOW (48), WEIS

(49), CREON (50), CASCON (51), SHERFACS (52), Real-Time Phoenix (23), and COEE (53) (see histories in (54) and (55)).

^{||}See (53) for a recent review of ontological depth and availability of Gold Standard example text.

Table 1. Ontological coverage of ICBe versus existing State of the Art

		Concept	Events Datasets						Episodes Datasets			
Domain	Player		ICBe	Phoenix	Terrier	ICCEWs	MIDs	Incidents	UCDP-GED	ICB	MIDs	COW
Domain	Players	Start (1–4) End (5, 6) N Coders Corpus Date Location	1918 2017 15k	1945 2019 8.5M	1977 2018 28.4M	1995 2020 17.5M	1993 2010 9.6K	1989 2015 128k		1918 2017 1K	1816 2014 5.9K	1816 2007 1K
Domain	Pieces	States (7–10) Subnational Actors (11–14) IGO/NGO (15–18) Civilians (19, 20)	Hand	Automated (CAMEO)			Hand	Hand		Hand	Hand	Hand
Domain	Pieces	Fatalities (21, 22) Force Size (23–25) Force Domain (26–29) Geography (location, territorial change) (30)	ICB Event	News Article		Mix Event	News Article	Mix Actor	News Event	Mix Event	Mix Event	Mix Actor
Domain	Think	Alert (Start/End Crisis) (31) Wishes (Desire/Fear) (32, 33) Evaluation (Victory/Defeat) (34) Aims (Territory, Policy, Regime, Preemption) (35) Awareness (Discover, Become Convinced) (36–38)										
Domain	Say	React to past event (Praise, Disapprove, Accept, Reject, Accuse) (39–41) Request future event (Appeal, Demand) (42) Predict future event (Promise, Threaten, Express Intent, Offer Without Condition) (43, 44) Predict with condition (Offer, Ultimatum) (45)	✓	✓	✓	✓						
Domain	Unarmed	Government (Leadership/Institution Change, Coup, Assassination) (46–50) By Civilians (Protest/Riot/Strike) (51) Against Civilians (Terrorism, Domestic Rights, Mass Killing, Evacuate) (52, 53) Diplomacy (Discussion, Meeting, Mediation, Break off negotiations, Withdraw/Expel Diplomats, Propoganda) (54) Legal Agreements (Sign Agreement, Settle Dispute, Join War on Behalf of, Ally, Mutual Defense Pact, Open Border, Cede Territory, Allow Inspections, Political Succession, Leave Alliance, Terminate Treaty) (55–57) Violate Agreement (Violate Terms of Agreement) (58) Mutual Cooperation or Directed Aid (Economic cooperation or Aid, Military Cooperation, Intelligence Cooperation, Unspecified) (59) Directed Aid (General Political Support, Economic Aid, Humanitarian Aid, Military Aid, Intelligence Aid, Unspecified Aid) (60, 61)	✓	✓	✓	✓						
Domain	Armed	Preparation (Alert, Mobilization, Fortify, Exercise, Weapons Test) (62) Maneuver (Deployment, Show of Force, Blockade, No Fly Zone, Border Violation) (63) Combat (Battle/Clash, Attack, Invasion/Occupation, Bombard, Cease Fire, Retreat) (64, 65) Strategic (Declare War, Join War, Continue Fighting, Surrender, End War, Withdraw from War, Switch Sides) (66, 67) Autonomy (Assert Political Control Over, Assert Autonomy Against, Annex, Reduce Control Over, Decolonize) (68–70)	✓	✓	✓	✓						

ICBe Coding Process

The ICBe ontology follows a hierarchical design philosophy where a smaller number of significant decisions are made early on and then progressively refined into more specific details (59).^{**} Each coder was instructed to first thoroughly read the full crisis narrative and then presented with a custom graphical user interface (SI Appendix, Fig. B1). Coders then proceeded sentence by sentence, choosing the number of events (0-3) that occurred, the highest behavior (thought, speech, or activity), a set of players (P), whether the means were primarily armed or unarmed, whether there was an increase or decrease in aggression (uncooperative/escalating or cooperative/de-escalating), and finally one or more non-mutually exclusive specific activities. Some additional details like location and timing information was always collected while other details were only collected if appropriate, e.g. force size, fatalities, domains, units, etc. A unique feature of the ontology is that thought, speech, and do behaviors can be nested into combinations, e.g. an offer for the U.S.S.R. to remove missiles from Cuba in exchange for the U.S. removing missiles from Turkey. Through compounding, the ontology can capture what players were said to have known, learned, or said about other specific fully described actions.

Each crisis was typically assigned to 2 expert coders and 2 novice coders with an additional tie-breaking expert coder assigned to sentences with high disagreement.^{††} For the purposes of measuring intercoder agreement and consensus, we temporarily disaggregate the unit of analysis to the Coder-Crisis-Sentence-Tag ($n=971,547$), where a tag is any unique piece of information a coder can associate with a sentence such as an actor, date, behavior, etc. We then aggregate those tags into final events ($n=18,309$), using a consensus procedure (SI Appendix, Algorithm B2) that requires a tag to have been chosen by at least one expert coder and either a majority of expert or novice coders. This screens noisy tags that no expert considered possible but leverages novice knowledge to tie-break between equally plausible tags chosen by experts.

Internal Consistency

We evaluate the internal validity of the coding process in several ways. For every tag applied we calculate the observed intercoder agreement as the percent of other coders who also applied that same tag (SI Appendix, Fig. B3). Across all concepts, the Top 1 Tag Agreement was low among novices (34%), moderate for experts (65%), and high (87%) following the consensus screening procedure.

We attribute the remaining disagreement primarily to three sources. First, we required coders to rate their confidence which was observed to be low for 20% of sentences- half due to a mismatch between the ontology and the text (“survey doesn’t fit event”-45%) and half due to a lack of information or confused writing in the source text (“more knowledge needed”-40%, “confusing sentence”-6%). Observed disagreement varied predictably with self reported confidence (SI Ap-

pendix, Fig. B4). Second, as intended agreement is higher (75-80%) for questions with fewer options near the root of the ontology compared to agreement for questions near the leafs of the ontology (50%-60%). Third, individual coders exhibiting nontrivial coding styles, e.g. some more expressive applying many tags per concept while others focused on only the single best match. We further observed unintended synonymy, e.g. the same information can be framed as either a threat to do something or a promise not to do something.

Case Study Selection

The remaining two qualities we seek to measure are recall and precision of ICBe events in absolute terms and relative to other existing systems. We provide full ICB narratives, ICBe coding in an easy to read iconographic form, and a wide range of visualizations for every case on the companion website. In this paper, we focus on two deep case studies. The first is the Cuban Missile Crisis (Figure 1) which took place primarily in the second half of 1962, involved the United States, the Soviet Union, and Cuba, and is widely known for bringing the world to the brink of nuclear war (hereafter Cuban Missiles). The second is the Crimea-Donbas Crisis (SI Appendix Figure D1) which took place primarily in 2014, involved Russia, Ukraine, and NATO, and within a decade spiraled into a full scale invasion (hereafter Crimea-Donbas). Both cases involve a superpower in crisis with a neighbor, initiated by a change from a friendly to hostile regime, with implications for economic and military security for the superpower, risked full scale invasion, and eventually invited intervention by opposing superpowers. We choose these cases because they are substantively significant to 20th and 21st century international relations, widely known across scientific disciplines and popular culture, and are sufficiently brief to evaluate in depth.

Recall

Recall measures the share of desired information recovered by a sequence of coded events, $Pr(E|H)$, and is poorly defined for historical episodes. First, there is no genuine ground truth about what occurred, only surviving texts about it. Second, there is no *a priori* guide to what information is necessary detail and what is ignorable trivia. History suffers from what is known as the Coastline Paradox (60) — it has a fractal dimension greater than one such that the more you zoom in the more detail you will find about individual events and in between every two discrete events. The ICBe ontology is a proposal about what information is important, but we need an independent benchmark to evaluate whether that proposal is a good one and that allows for comparing proposals from event projects that had different goals. We need a yardstick for history.

Our strategy for dealing with both problems is a plausibly objective yardstick called a synthetic historical narrative. For both case studies, we collect a large diverse corpus of narratives spanning timelines, encyclopedia entries, journal articles, news reports, websites, and government documents. Using natural language processing (fully described in SI Appendix, Algorithm C1), we identify details that appear across multiple accounts. The more accounts that mention a detail, the more central it is to understanding the true historical episode. The theoretical motivation is that authors face word limits

^{**}This process quickly focuses the coder on a smaller number of relevant options while also allowing them to apply multiple tags if the sentence explicitly includes more than one or there is insufficient evidence to choose only one tag. The guided coding process also allows for the possibility that earlier coarse decisions have less error than later fine-grained decisions.

^{††}Expert coders were graduate students or postgraduates who collaboratively developed the ontology and documentation for the codebook. Undergraduate coders were students who engaged in classroom workshops.

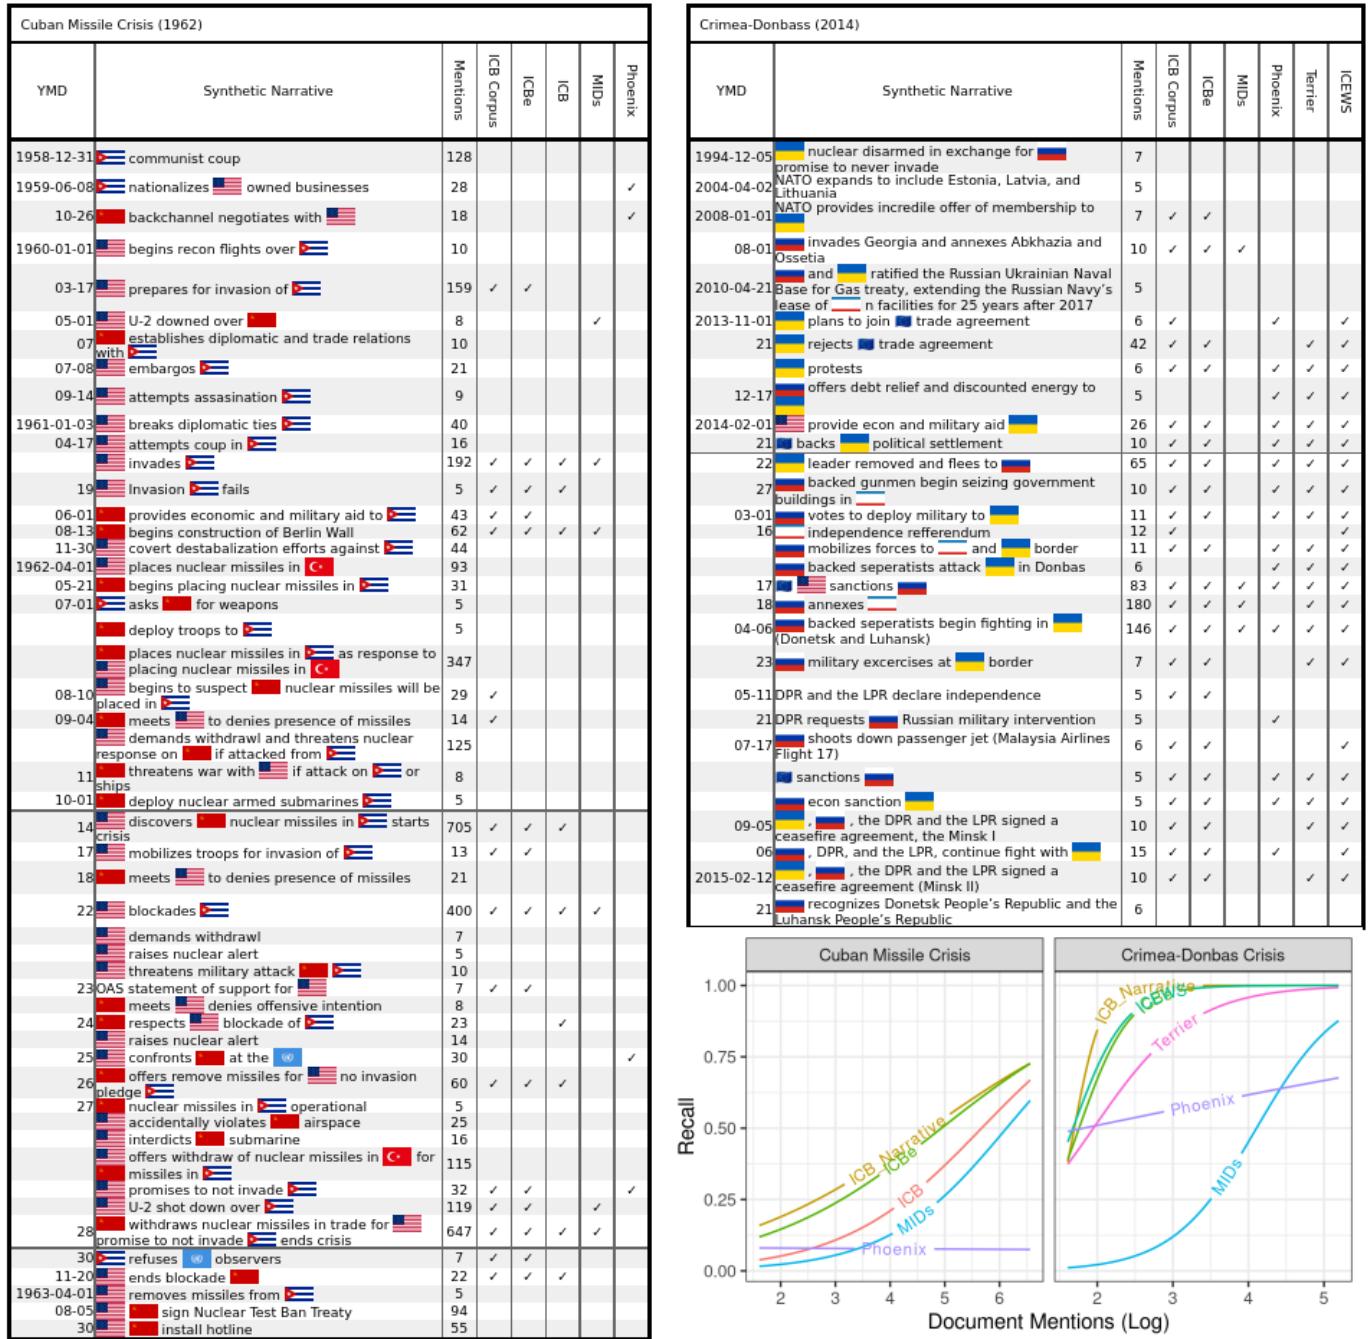
which force them to pick and choose which details to include, and they choose details which serve the specific context of the document they are producing. With a sufficiently large and diverse corpus of documents, we can vary the context while holding the overall episode constant and see which details tend to be invariant to context. Intuitively, a high quality event dataset should have high recall for context invariant details both because of their broader relevance and also because they are easier to find in source material.

Synthetic historical narratives for Cuban Missiles (51 events drawn from 2020 documents) and Crimea-Donbas (30 events drawn from 971 documents) appear in Figure 2. Each row represents a detail which appeared in at least five documents along with an approximate start date, a hand written summary, the number of documents it was mentioned in, and whether it could be identified in the text of our ICB corpus, in our ICBe events, or any of the competing systems.

From them, we draw several stylized facts. First, there is substantial variation in which details any one document will choose to include. Our ground truth ICB narratives included 17/51 and 23/30 of the events from the synthetic narrative, while including other details that are not in the synthetic narrative. Second, mentions of a detail across accounts is exponentially distributed with context invariant details appearing dozens to hundreds of times more than context dependent details. Third, crisis start and stop dates are arbitrary and the historical record points to many precursor events as necessary detail for understanding later events, e.g. the U.S. was in a *de facto* grey scale war with Cuba before it invited Soviet military protection (61) and Ukraine provided several security guarantees to Russia that were potentially undone, e.g. a long term lease on naval facilities in Crimea. Fourth, we find variation between the two cases. Cuban Missiles has a cleaner canonical end with the Soviets agreeing to withdraw missiles while Crimea-Donbas meekly ends with a second cease fire agreement (Minsk II) but continued fighting. The canonical narrative of Cuban Missile also includes high level previously classified details, while the more recent Crimea-Donbas case reflects primarily public reporting.

We find substantive variation in recall across systems. Recall for each increases in the number of document mentions which is an important sign of validity for both them and our benchmark. The one outlier is Phoenix which is so noisy that it's flat to decreasing in mentions. The two episode level datasets have very low coverage of contextual details. The two other dictionary systems ICEWs and Terrier have high coverage, with ICEWs outperforming Terrier. ICBe strictly dominates all of the systems but ICEWs in recall though we note that the small sample sizes mean these systems should be considered statistically indistinguishable. Importantly our corpus of ICB narratives has very high recall of frequently mentioned details giving us confidence in how those summaries were constructed, and ICBe lags only slightly behind showing that it left very little additional information on the table.

Fig. 2. Measuring Recall with Synthetic Historical Narratives

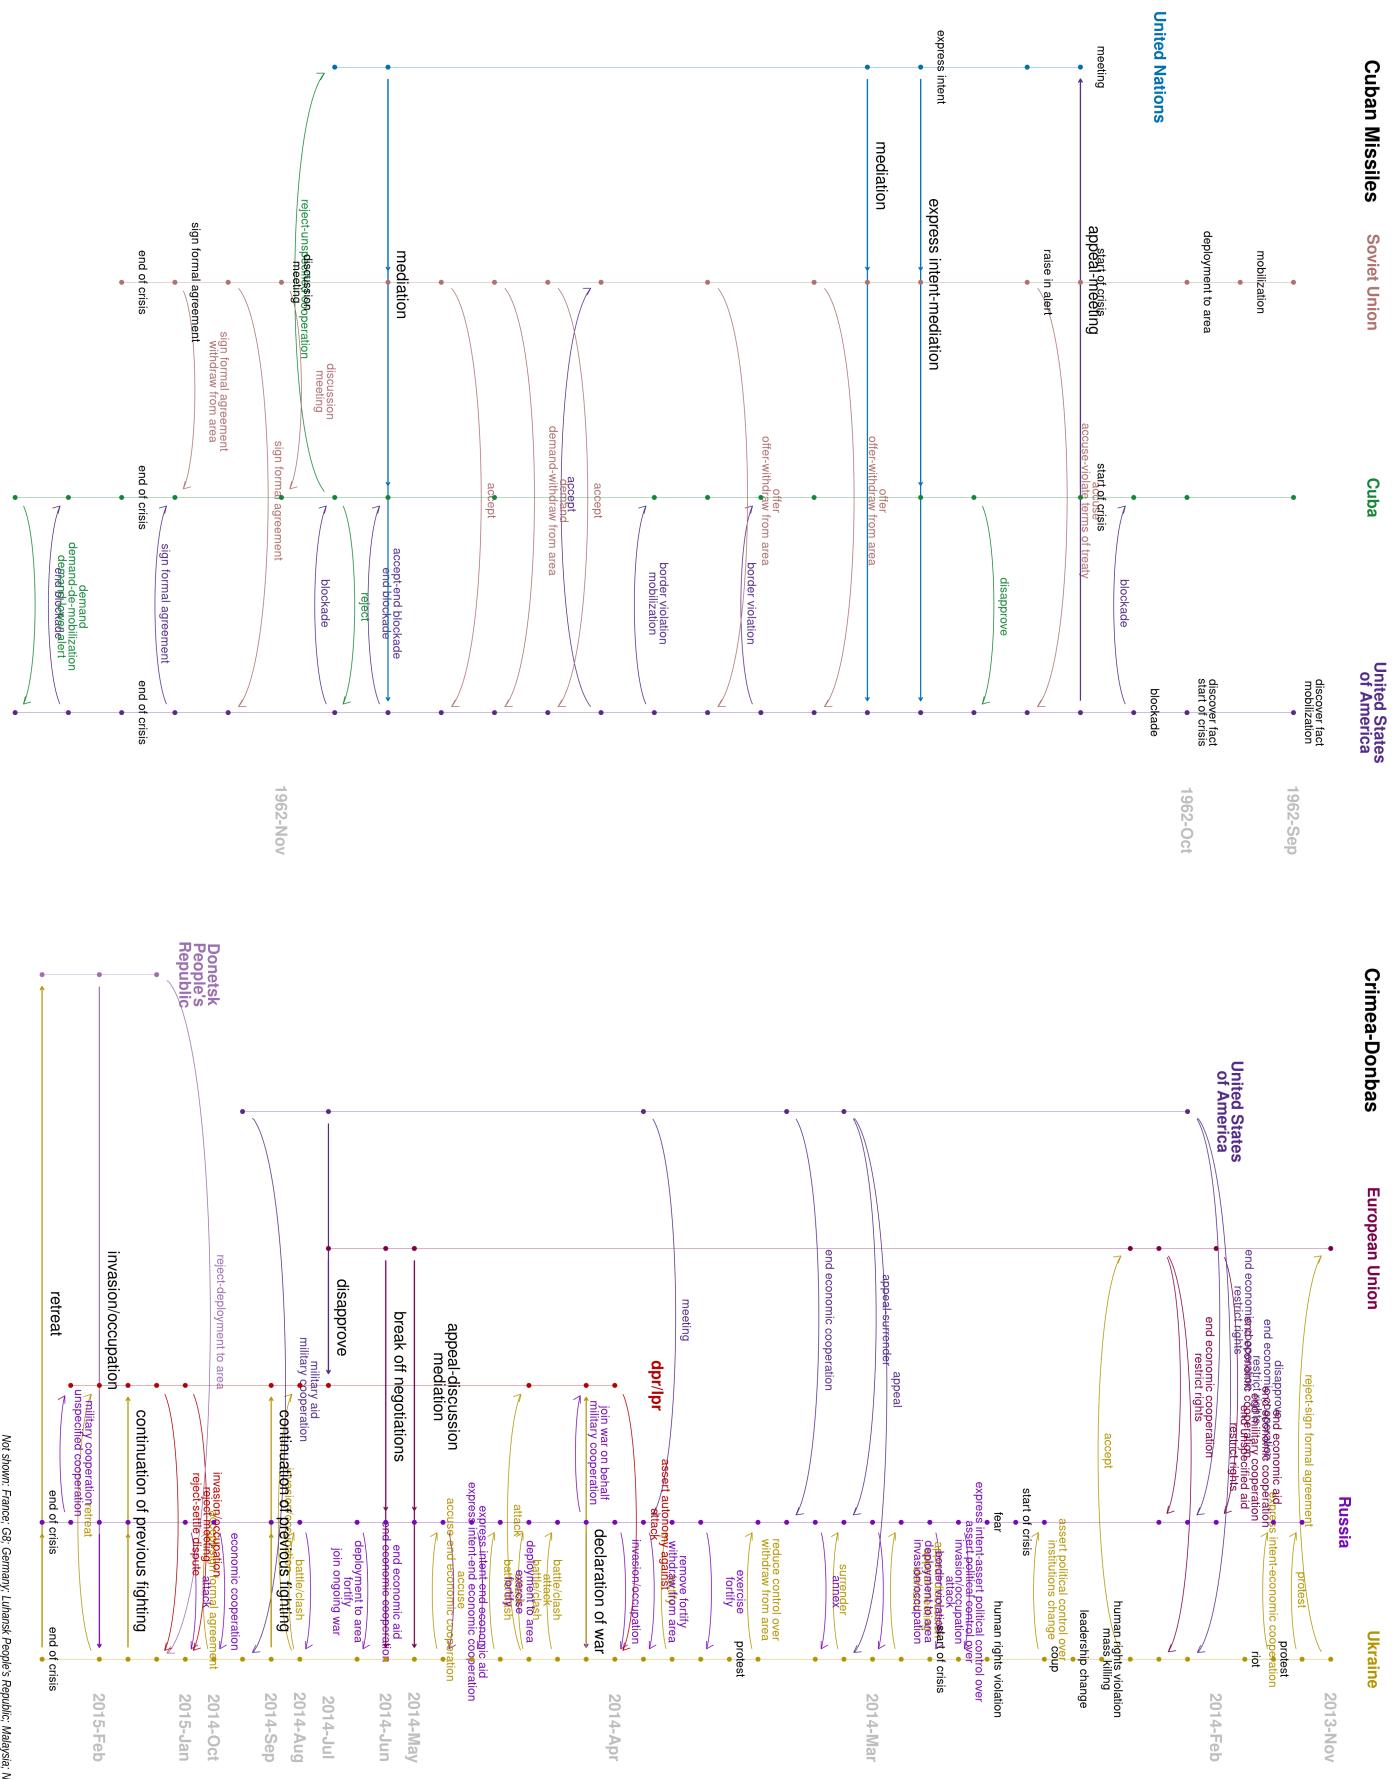


Notes: Synthetic narratives combine several thousand accounts of each crisis into a single timeline of events, taking only those mentioned in at least 5 or more documents. Checkmarks represent whether that event could be hand matched to any detail in the ICB corpus, ICBe dataset, or any of the other event datasets.

Precision

The other side of event measurement is precision, the degree to which a sequence of events correctly and usefully describes the information in history, $Pr(H|E)$. It does little good to recall a historical event but too vaguely (e.g. MIDs describes the Cuban Missile crisis as a blockade, a show of force, and a stalemate) or with too much error (e.g. ICEWS records 263 “Detonate Nuclear Weapons” events between 1995-2019) to be useful for downstream applications. ICBe’s ontology and coding system is designed to strike a balance so that the most important information is recovered accurately but also abstracted to a level that is still useful and interpretable. You should be able to lay out events of a crisis on a timeline, as in Figure 3, and read off the macro structure of an episode from each individual move. We call this visualization a crisis map, a directed graph intersected with a timeline, and provide crisis maps for every event dataset for each case study (SI Appendix, Fig. D3 and D4) and all crises on the companion website.

Fig. 3. Crisis Maps



Not shown: Organization of American States; Warsaw Pact

Not shown: France, GA, Germany, Luhansk People's Republic, Malaysia, MATO, Organization for Security and Cooperation in Europe, Trident Covert Group, UN, Ukraine, United Nations

We further want to verify individual event codings, which we can do in the case of ICBe because each event is mapped to a specific span of text. We develop the iconography system for presenting event codings as coherent statements that can be compared side by side to the original source narrative as for Cuban Missiles (Figure 1), Crimea-Donbas (SI Appendix Table D1), and for every case on the companion website. We further provide a stratified sample of event codings alongside their source text (SI Appendix Table D2).

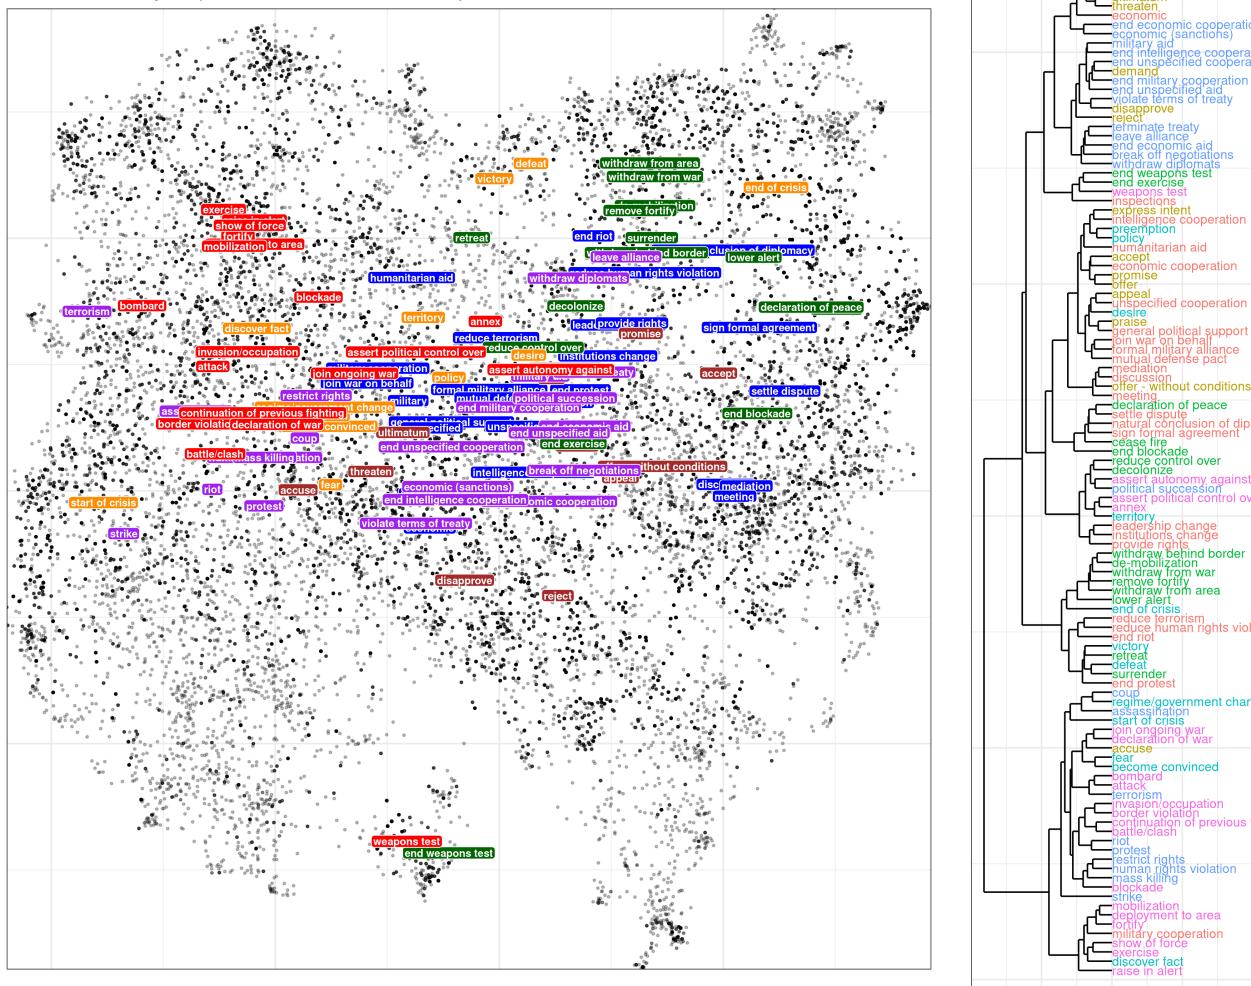
We find both the visualizations of macro structure and head-to-head comparisons of ICBe codings to the raw text to strongly support the quality of ICBe, but as with recall we seek a more objective detached universal benchmark. Our proposed measure is a reconstruction task to see whether our intended ontology can be recovered through only unsupervised clustering of sentences they were applied to. Figure 4 shows the location of every sentence from the ICBe corpus in semantic space as embedded using the same large language model as before, and the median location of each ICBe event tag applied to those sentences.^{††} Labels reflect the individual leaves of the ontology and colors reflect the higher level coerce branch nodes of the ontology. If ICBe has high precision, substantively similar tags ought to have been applied to substantively similar source text, which is what we see both in two dimensions in the main plot and via hierarchical clustering on all dimensions in the dendrogram along the righthand side.^{§§}

^{††}We preprocess sentences to replace named entities with a generic Entity token.

^{§§}Hierarchical clustering on cosine similarity and with Ward's method.

Fig. 4. ICB event codings in comparison to Semantic Embeddings from source sentences

Variation in Tags by Semantic Embeddings of Source Sentences
Sentence Embedding Paraphrase-MPNET-base-v2, UMAP Projection



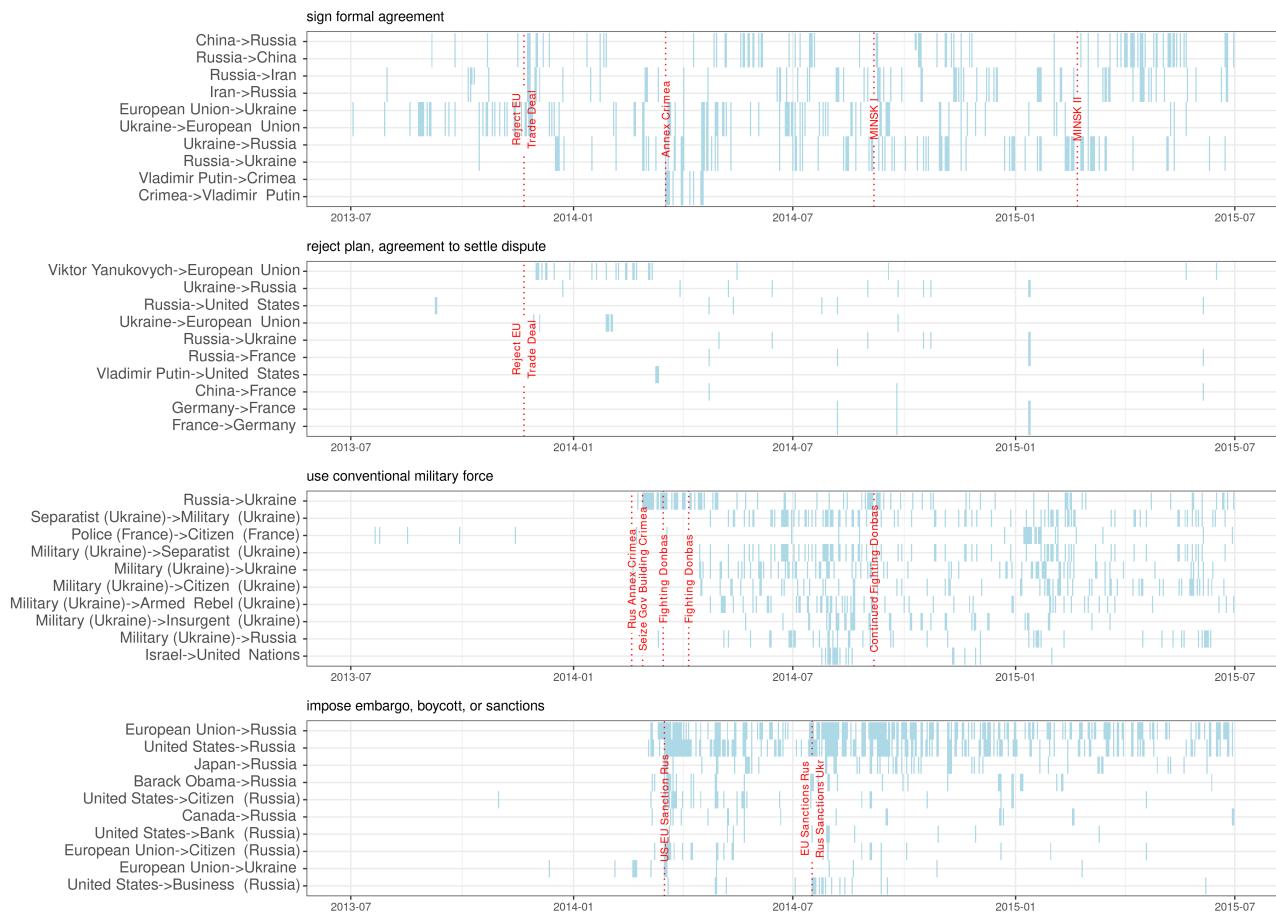
Notes: Dots represent individual ICB narrative sentences, as embeded by the Paraphrase-MPNET-base-v2 large language model and flattened into two dimensions with UMAP. Text labels reflect individual leaves of the ICB ontology, and colors represent intermediate branches of the ontology. Label placement is the median of all of the sentences that tag was applied to by the coders. The dendrogram shows hierarchical clustering of the tags. If ICB precision is high, the sentences tags were applied to ought to say similar things, and the intended shape of the ontology ought to visually recognizable.

Finally, how does ICBe’s precision compare to the existing state of the art? The crisis-maps reveal the episode level datasets like MIDs or the original ICB are too sparse and vague to reconstruct the structure of the crisis. On the other end of the spectrum, the high recall dictionary based event datasets like Terrier and ICEWs produce so many noisy events (several hundreds thousands) that even with heavy filtering their crisis maps are completely unintelligible. Further, because of copyright issues, none of these datasets directly provide the original text spans making event level precision difficult to verify.

However, given their high recall on our task and the global and real-time coverage of dictionary based event systems, we want to take seriously the possibility that some functional transformation could recover the precision of ICBe. For example, (62) attempts to correct for the mechanically increasing amount of news coverage each year by detrending violent event counts from Phoenix using a human coded baseline. Others have focused on verifying precision for ICEWs on specific subsets of details against known ground truths, e.g. geolocation (63), protest events (80%) (64), anti-government protest networks (46.1%) (65).

We take the same approach here in Figure 5, selecting four specific CAMEO event codings and checking how often they reflect a true real world event. We choose four event types around key moments in the crisis. The start of the crisis revolves around the Ukraine backing out of trade deal with the EU in favor of Russia, but the “sign formal agreement” acts more like a topic detector with dozens of events generated by discussions of a possible agreement but not the actual agreement which never materialized. The switch is caught by the “reject plan, agreement to settle dispute”, but also continues for Victor Yanukovych for even after he was removed from power because of articles retroactively discussing the cause of his removal. Events for “use conventional military force” capture a threshold around the start of hostilities and who the participants were but not any particular battles or campaigns. Likewise, “impose embargo, boycott, or sanctions” captures the start of waves of sanctions and from who but are effectively constantly as the news coverage does not distinguish between subtle changes or additions. In sum, dictionary based methods on news corpora tend to have high recall because they parse everything in the news, but for the same reason their specificity for most event types is too low to back out individual chess like sequencing that ICBe aims to record.

Fig. 5. ICEWs Counts of Events per Day by Type during the Crimea-Donbas Crisis



Notes: Unit of analysis is the Dyad-Day. Edges <-> indicates undirected dyad and -> indicates directed dyad. Top 10 most active dyads per category shown. Red text shows events from the synthetic narrative relative to that event category. Event types are grouped by hand along similar themes. Blue bars indicate an event recorded by ICEWs for that dyad on that day.

Conclusion

We investigated event extraction from narratives describing key historical episodes in international relations. We synthesized a prior belief about the latent unobserved phenomena that drive these events in international relations and proposed a mapping to observable concepts that enter into the observed historical record. We designed an ontology with high coverage over those concepts and developed a training procedure and technical stack for human coding of historical texts. Multiple validity checks find the resulting codings have high internal validity (e.g. intercoder agreement) and external validity (i.e. matching source material in both micro-details at the sentence level and macro-details spanning full historical episodes). Further, these codings perform much better in terms of recall, precision, coverage, and overall coherence in capturing these historical episodes than existing event systems used in international relations.

We release several open-source products along with supporting code and documentation to further advance the study of IR, event extraction, and natural language processing. The first is the International Crisis Behavior Events (ICBe) dataset, an event-level aggregation of what took place during the crises identified by the ICB project. These data are appropriate for statistical analysis of hard questions about the sequencing of events (e.g. escalation and de-escalation of conflicts). Second, we provide a coder-level disaggregation with multiple codings of each sentence by experts and undergrads that allows for the introduction of uncertainty and human interpretation of events. Further, we release a direct mapping from the codings to the source text at the sentence level as a new resource for natural language processing. Finally, we provide a companion website that incorporates detailed visualizations of all of the data introduced here (www.crisisevents.org).

ACKNOWLEDGMENTS. We thank the ICB Project and its directors and contributors for their foundational work and their help with this effort. We make special acknowledgment of Michael Brecher for helping found the ICB project in 1975, creating a resource that continues to spark new insights to this day. We thank the many undergraduate coders for their patience and dedication. Thanks to the Center for Peace and Security Studies and its membership for comments. Special thanks to Rebecca Cordell, Philip Schrot, Zachary Steinert-Threlkeld, and Zhanna Terechshenko for generous feedback. Thank you to the cPASS research assistants that contributed to this project: Helen Chung, Daman Heer, Syeda ShahBano Ijaz, Anthony Limon, Erin Ling, Ari Michelson, Prithviraj Pahwa, Gianna Pedro, Tobias Stodiek, Yiyi ‘Effie’ Sun, Erin Werner, Lisa Yen, and Ruixuan Zhang. This project was supported by a grant from the Office of Naval Research [N00014-19-1-2491] and benefited from the Charles Koch Foundation’s support for the Center for Peace and Security Studies.

1. Brecher M (1999) International studies in the twentieth century and beyond: Flawed dichotomies, synthesis, cumulation: ISA presidential address. *International Studies Quarterly* 43(2):213–264.
2. Reiter D (2015) *Should We Leave Behind the Subfield of International Relations?* *Annu Rev Polit Sci* 18(1):481–499.
3. Brecher M, Wilkenfeld J (1982) *Crises in World Politics*. *World Politics* 34(3):380–417.
4. Beardsley K, James P, Wilkenfeld J, Brecher M (2020) The International Crisis Behavior Project. doi:[10.1093/acrefore/9780190228637.013.1638](https://doi.org/10.1093/acrefore/9780190228637.013.1638).

5. Palmer G, et al. (2021) *The MID5 Dataset, 2011–2014: Procedures, coding rules, and description*. *Conflict Management and Peace Science*:0738894221995743.
6. Gibler DM (2018) *International Conflicts, 1816–2010: Militarized Interstate Dispute Narratives* (Rowman & Littlefield) Available at: https://books.google.com?id=_4VTDwAAQBAJ.
7. Maoz Z, Johnson PL, Kaplan J, Ogunkoya F, Shreve AP (2019) *The Dyadic Militarized Interstate Disputes (MIDs) Dataset Version 3.0: Logic, Characteristics, and Comparisons to Alternative Datasets*. *Journal of Conflict Resolution* 63(3):811–835.
8. Sarkees MR, Wayman F (2010) *Resort to war: 1816–2007* (CQ Press).
9. Reiter D, Stam AC, Horowitz MC (2016) *A Revised Look at Interstate Wars, 1816–2007*. *Journal of Conflict Resolution* 60(5):956–976.
10. Ralph Sundberg, Mihai Croicu (2016) *UCDP GED Codebook version 5.0* (Department of Peace and Conflict Research, Uppsala University).
11. Pettersson T, Eck K (2018) *Organized violence, 1989–2017*. *Journal of Peace Research* 55(4):535–547.
12. Raleigh C, Linke A, Hegre H, Karlsen J (2010) Introducing ACLED: An armed conflict location and event dataset: Special data feature. *Journal of peace research* 47(5):651–660.
13. Felbermayr G, Kirilakha A, Syropoulos C, Yalcin E, Yotov YV (2020) *The global sanctions data base*. *European Economic Review* 129:103561.
14. Barari S, Kim IS Democracy and Trade Policy at the Product Level: Evidence from a New Tariff-line Dataset. 16.
15. Kinne BJ (2020) *The Defense Cooperation Agreement Dataset (DCAD)*. *Journal of Conflict Resolution* 64(4):729–755.
16. Owsiak AP, Cuttner AK, Buck B (2018) *The International Border Agreements Dataset*. *Conflict Management and Peace Science* 35(5):559–576.
17. Vabulas F, Snidal D (2021) *Cooperation under autonomy: Building and analyzing the Informal Intergovernmental Organizations 2.0 dataset*. *Journal of Peace Research* 58(4):859–869.
18. Frederick BA, Hensel PR, Macaulay C (2017) *The Issue Correlates of War Territorial Claims Data, 1816–2001*. *Journal of Peace Research* 54(1):99–108.
19. Moyer JD, Turner SD, Meisel CJ (2020) *What are the drivers of diplomacy? Introducing and testing new annual dyadic data measuring diplomatic exchange*. *Journal of Peace Research*:0022343320929740.
20. Sechser TS (2011) *Militarized Compellent Threats, 1918–2001*. *Conflict Management and Peace Science* 28(4):377–401.
21. Li Q, et al. (2021) A Comprehensive Survey on Schema-based Event Extraction with Deep Learning. Available at: <http://arxiv.org/abs/2107.02126> [Accessed September 10, 2021].
22. Halterman A (2020) Extracting Political Events from Text Using Syntax and Semantics.
23. Brandt PT, D’Orazio V, Holmes J, Khan L, Ng V (2018) Phoenix Real-Time Event Data. Available at: <http://eventdata.utdallas.edu>.
24. Boschee E, et al. (2015) ICEWS coded event data. *Harvard Dataverse* 12.
25. Hegre H, Croicu M, Eck K, Höglbladh S (2020) Introducing the UCDP Candidate Events Dataset. *Research & Politics* 7(3):2053168020935257.
26. Grant C, Halterman A, Irvine J, Liang Y, Jabr K (2017) OU Event Data Project. Available at: <https://osf.io/4m2u7/> [Accessed September 1, 2021].

27. Zhang H, Pan J (2019) CASM: A Deep-Learning Approach for Identifying Collective Action Events with Text and Image Data from Social Media. *Sociological Methodology* 49(1):1–57.
28. Steinert-Threlkeld ZC (2019) *The Future of Event Data Is Images*. *Sociological Methodology* 49(1):68–75.
29. Brecher M, James P, Wilkenfeld J (2000) Crisis escalation to war: Findings from the International Crisis Behavior Project. *What Do We Know About War*.
30. Wilkenfeld J, Brecher M (2000) Interstate crises and violence: Twentieth-century findings. *Handbook of war studies II*:282–300.
31. James P (2019) *What do we know about crisis, escalation and war? A visual assessment of the International Crisis Behavior Project*. *Conflict Management and Peace Science* 36(1):3–19.
32. Iakhnis E, James P (2019) *Near crises in world politics: A new dataset*. *Conflict Management and Peace Science*:0738894219855610.
33. Brecher M, Wilkenfeld J, Beardsley KC, James P, Quinn D (2017) *International Crisis Behavior Data Codebook* Available at: <http://sites.duke.edu/icbdata/data-collections/>.
34. Brecher M, Wilkenfeld J (1997) *A Study of Crisis* (University of Michigan Press).
35. Holsti OR (1965) *The 1914 Case*. *The American Political Science Review* 59(2):365–378.
36. Paige GD (1968) *The Korean Decision, June 24–30, 1950* (Free Press).
37. Allison GT, Zelikow P (1971) *Essence of decision: Explaining the Cuban missile crisis* (Little, Brown Boston).
38. Snyder GH, Diesing P (1977) *Conflict among nations: Bargaining and decision making in international crises* (Princeton University Press).
39. Gavin FJ (2014) *History, Security Studies, and the July Crisis*. *Journal of Strategic Studies* 37(2):319–331.
40. George AL, Smoke R (1974) *Deterrence in American foreign policy: Theory and practice* (Columbia University Press).
41. Gaddis JL (1987) *Expanding the Data Base: Historians, Political Scientists, and the Enrichment of Security Studies*. *International Security* 12(1):3–21.
42. Brecher M, James P (1988) Patterns of crisis management. *Journal of Conflict Resolution* 32(3):426–456.
43. Kang DC, Lin AY-T (2019) US bias in the study of Asian security: Using Europe to study Asia. *Journal of Global Security Studies* 4(3):393–401.
44. Hewitt JJ (2001) *Engaging International Data in the Classroom: Using the ICB Interactive Data Library to Teach Conflict and Crisis Analysis*. *Int Stud Perspect* 2(4):371–383.
45. Miller GA (1995) *WordNet: A lexical database for English*. *Commun ACM* 38(11):39–41.
46. Althaus S, Bajjalieh J, Carter JF, Peyton B, Shalmon DA (2019) Cline Center Historical Phoenix Event Data Variable Descriptions. *Cline Center Historical Phoenix Event Data*.
47. Sundberg R, Melander E (2013) Introducing the UCDP georeferenced event dataset. *Journal of Peace Research* 50(4):523–532.
48. Leng RJ, Singer JD (1988) *Militarized Interstate Crises: The BCOW Typology and Its Applications*. *International Studies Quarterly* 32(2):155–173.
49. McClelland C (1978) World event/interaction survey, 1966–1978. *WEIS Codebook ICPSR 5211*.
50. Hermann C (1984) Comparative Research on the Events of Nations (CREON) Project: Foreign Policy Events, 1959–1968: Version 1. doi:[10.3886/ICPSR05205.V1](https://doi.org/10.3886/ICPSR05205.V1).
51. Bloomfield LP, Moulton A (1989) CASCON III: Computer-aided system for analysis of local conflicts. *MIT Center for International Studies, Cambridge*.
52. Sherman FL (2000) SHERFACS: A Cross-Paradigm, Hierarchical, and Contextually-Sensitive International Conflict Dataset, 1937–1985: Version 1. doi:[10.3886/ICPSR02292.V1](https://doi.org/10.3886/ICPSR02292.V1).
53. Balali A, Asadpour M, Jafari SH (2021) COFFEE: A Comprehensive Ontology for Event Extraction from text. doi:[10.48550/arXiv.2107.10326](https://doi.org/10.48550/arXiv.2107.10326).
54. Merritt RL (1994) Measuring events for international political analysis. *International Interactions* 20(1–2):3–33.
55. Schrodt PA, Hall B (2006) Twenty years of the Kansas event data system project. *The political methodologist* 14(1):2–8.
56. Brecher M, Wilkenfeld J, Beardsley K, James P, Quinn D International Crisis Behavior Data Codebook, Version 12. 69.
57. Braithwaite A (2010) *MIDLOC: Introducing the Militarized Interstate Dispute Location dataset*. *Journal of Peace Research* 47(1):91–98.
58. Braithwaite A (2009) Codebook for the Militarized Interstate Dispute Location (MIDLOC) Data, v 1.0. *University College London*.
59. Brust C-A, Denzler J (2020) Integrating domain knowledge: Using hierarchies to improve deep classifiers. Available at: <http://arxiv.org/abs/1811.07125> [Accessed September 7, 2021].
60. Mandelbrot BB (1983) *The fractal geometry of nature* (Freeman, New York).
61. Cormac R, Aldrich RJ (2018) *Grey is the new black: Covert action and implausible deniability*. *International Affairs* 94(3):477–494.
62. Terechshenko Z (2020) *Hot under the collar: A latent measure of interstate hostility*. *Journal of Peace Research* 57(6):764–776.
63. Cook SJ, Weidmann NB (2019) *Lost in Aggregation: Improving Event Analysis with Report-Level Data*. *American Journal of Political Science* 63(1):250–264.
64. Wüest B, Lorenzini J (2020) External validation of protest event analysis. *Contention in times of crisis: Recession and political protest in thirty European countries*:49–78.
65. Jäger K The Limits of Studying Networks with Event Data: Evidence from the ICEWS Dataset.